

# CSCI 381/780: Parallel and Distributed Computing

## Parallel Programming with MPI



Jun Li

Queens College & Graduate Center

[jun.li@qc.cuny.edu](mailto:jun.li@qc.cuny.edu)

# Roadmap

**MPI Overview**

**Process Model & Language Bindings**

**Messages & Point-to-Point Communication**

**Nonblocking Communication**

**Collective Communication**

**Error Handling**

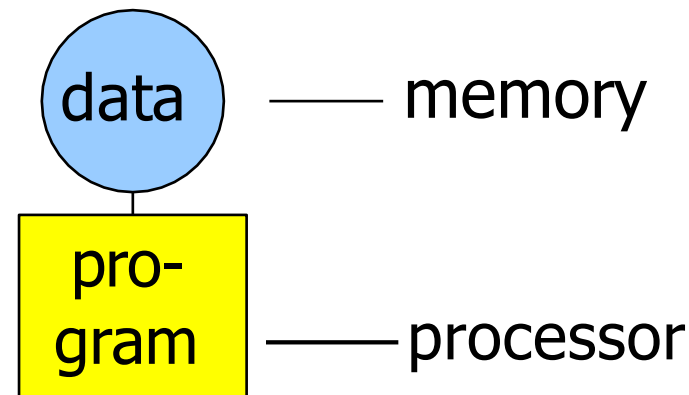
**Shared Memory**

**Parallel I/O**

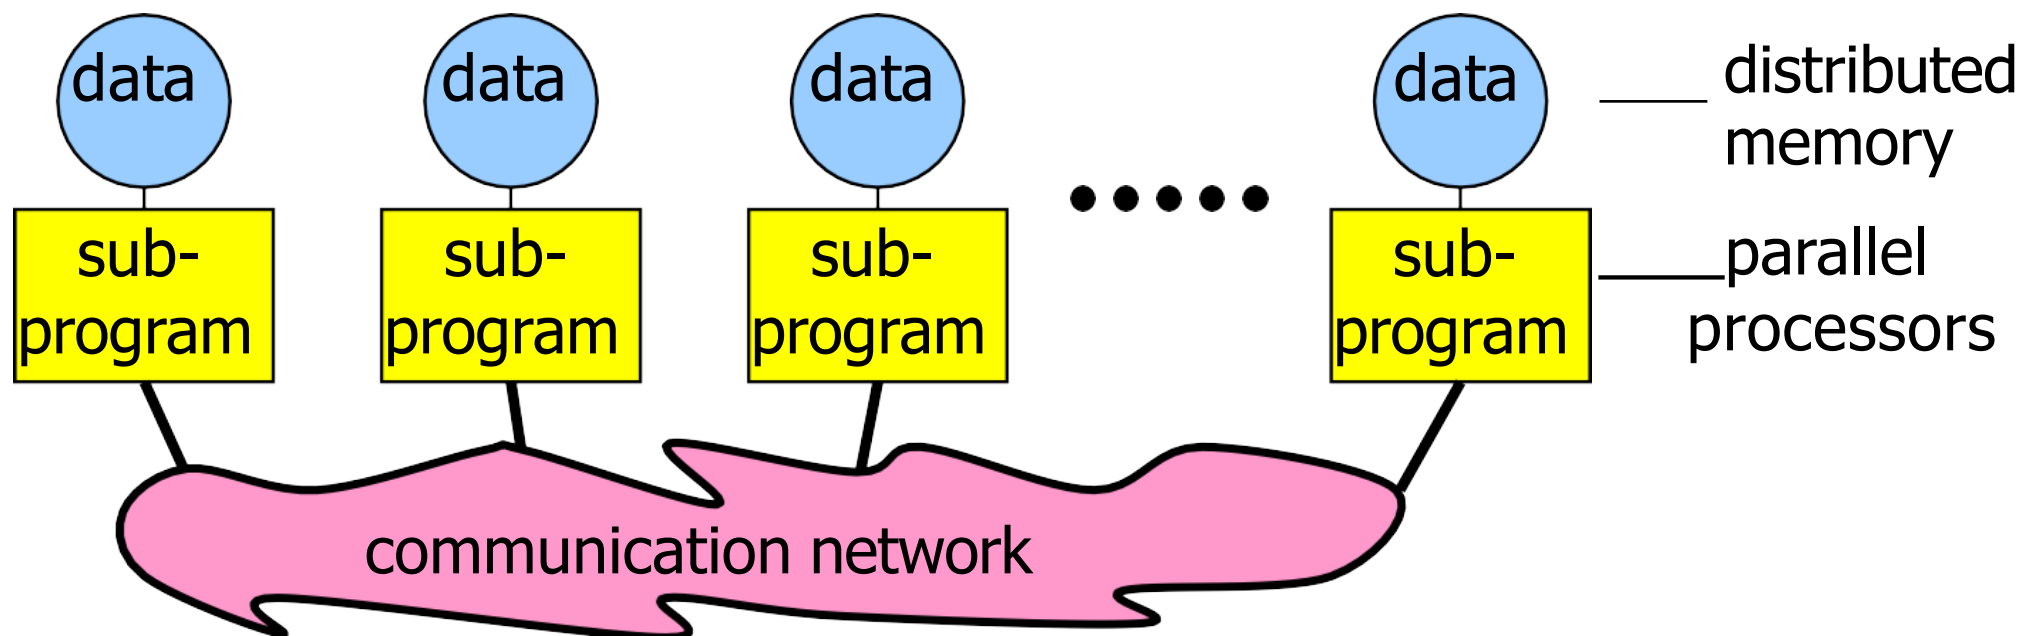
# MPI Overview

# The Message-Passing Programming Paradigm

- Sequential Programming Paradigm

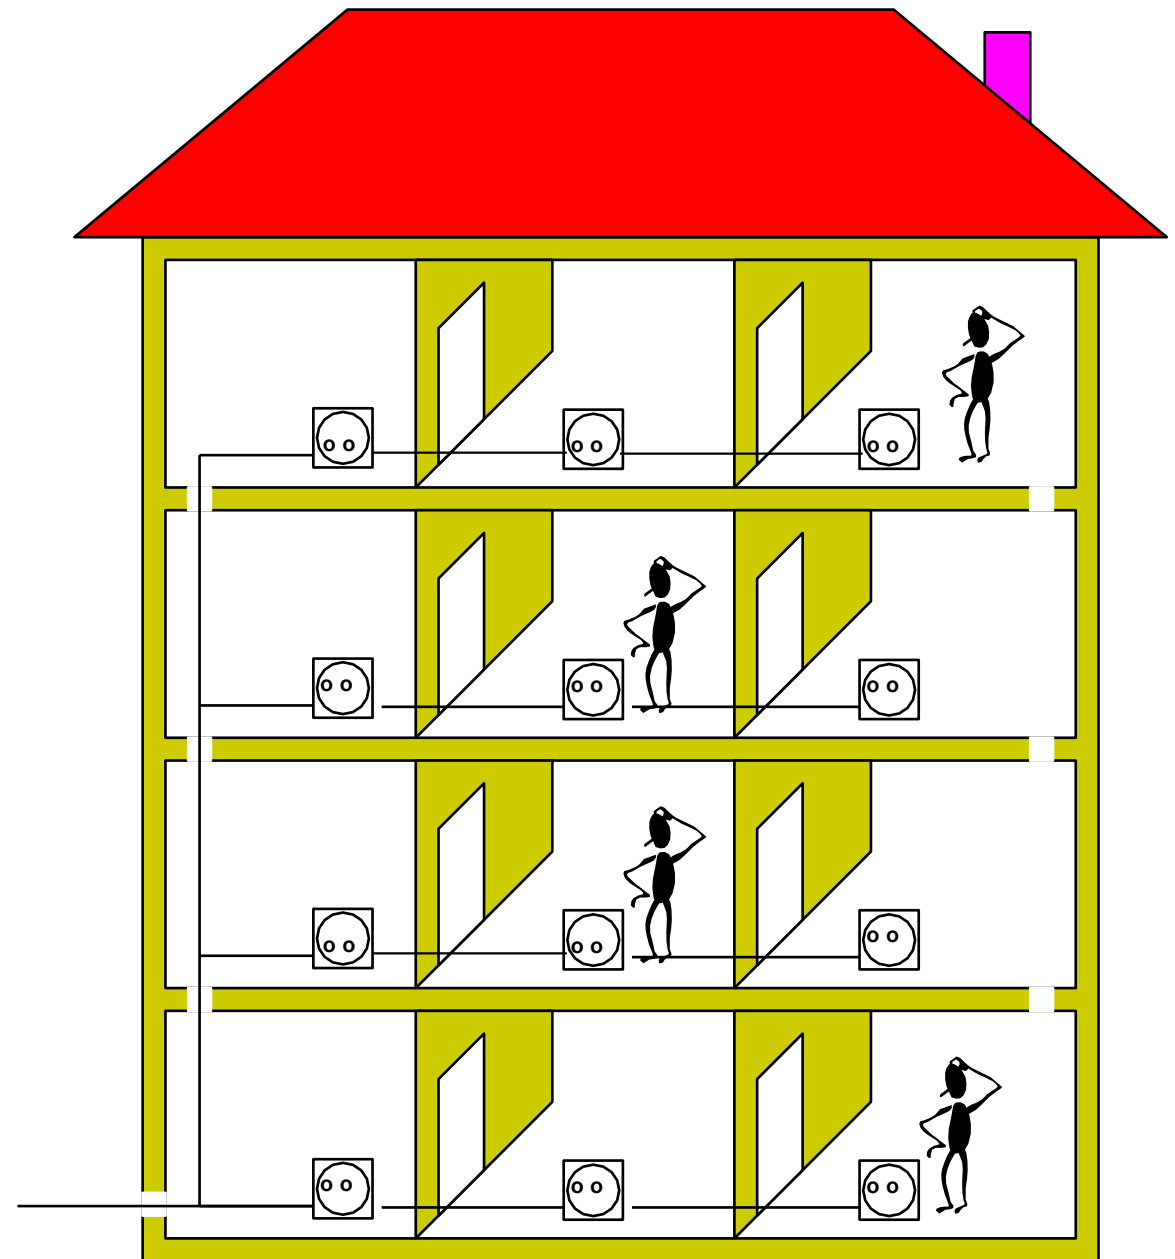


- Message-Passing Programming Paradigm




# Analogy: Electric Installations in Parallel

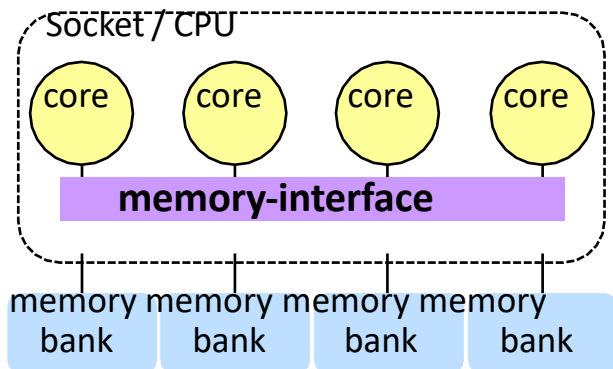
- MPI sub-program  
= work of one electrician  
on one floor
- data  
= the electric installation
- MPI communication  
= real communication  
to guarantee that the wires  
are coming at the same  
position through the floor



# Parallel hardware architectures



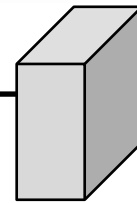
  
shared memory



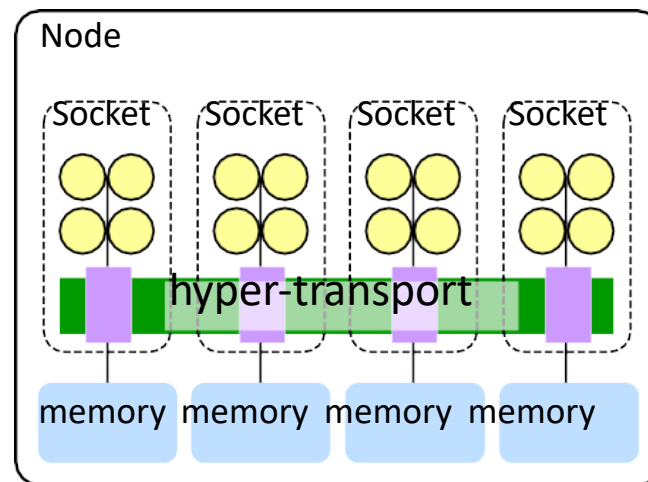
## Socket/CPU

→ **memory interface**

**UMA (uniform memory access) SMP (symmetric multi-processing)** All cores connected to all memory banks with same speed



distributed memory



## Node

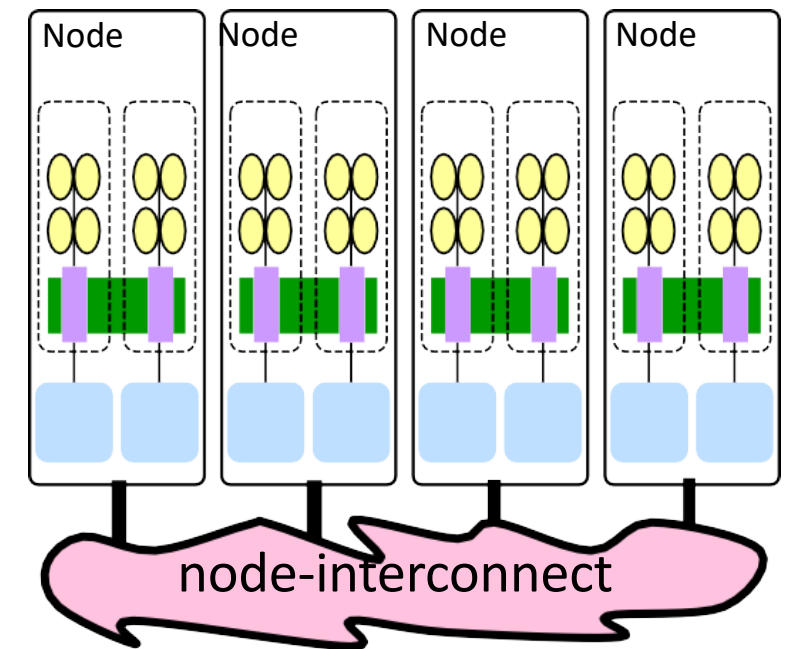
→ **hyper-transport**

**ccNUMA (cache-coherent non-uniform memory access)**

Shared memory programming is possible

**Performance problems:**

- Threads should be **pinned** to the physical sockets
- **First-touch** strategy is needed to minimize remote memory access



## Cluster

→ **node-interconnect**

**NUMA (non-uniform memory access)**

**!! fast access only on its own memory !!**

**Many programming options:**

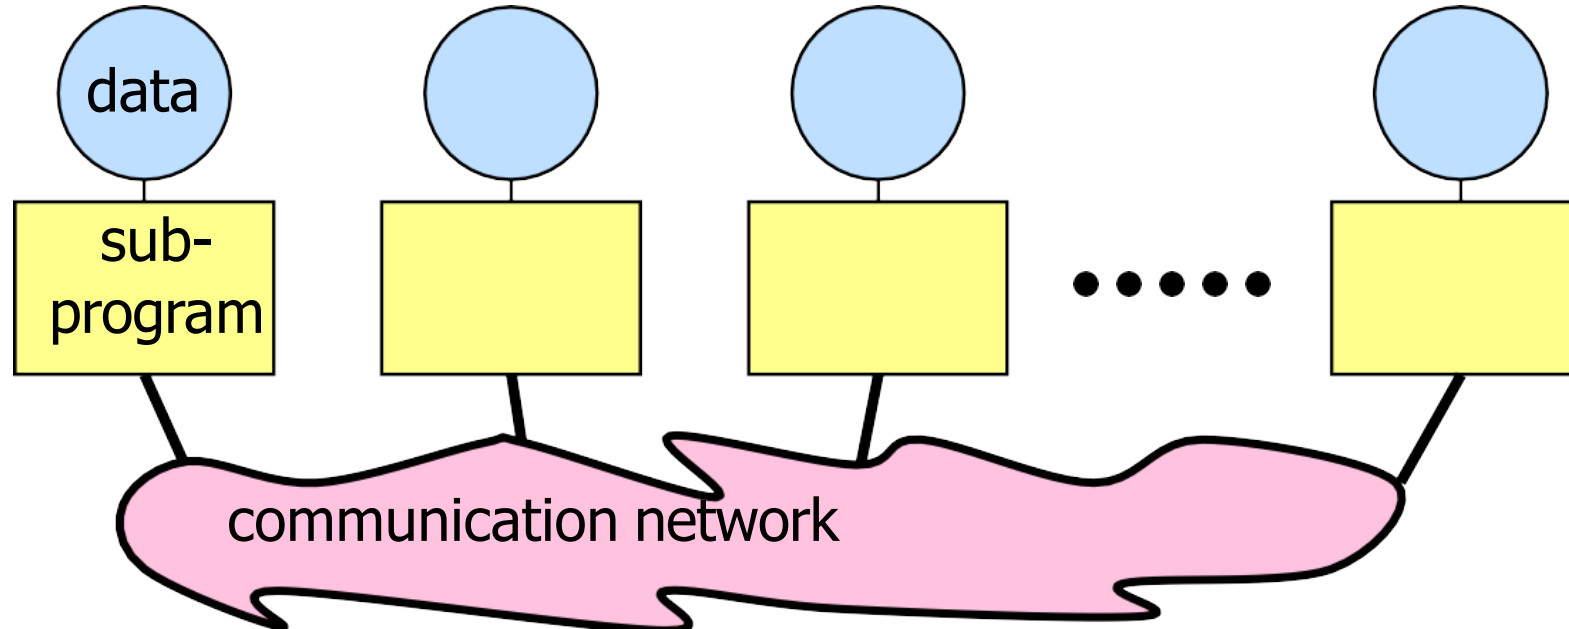
- Shared memory / symmetric multi-processing inside of each node
- distributed memory parallelization on the node interconnect
- **Or simply one MPI process on each core**

Shared memory programming with OpenMP

**MPI works everywhere**

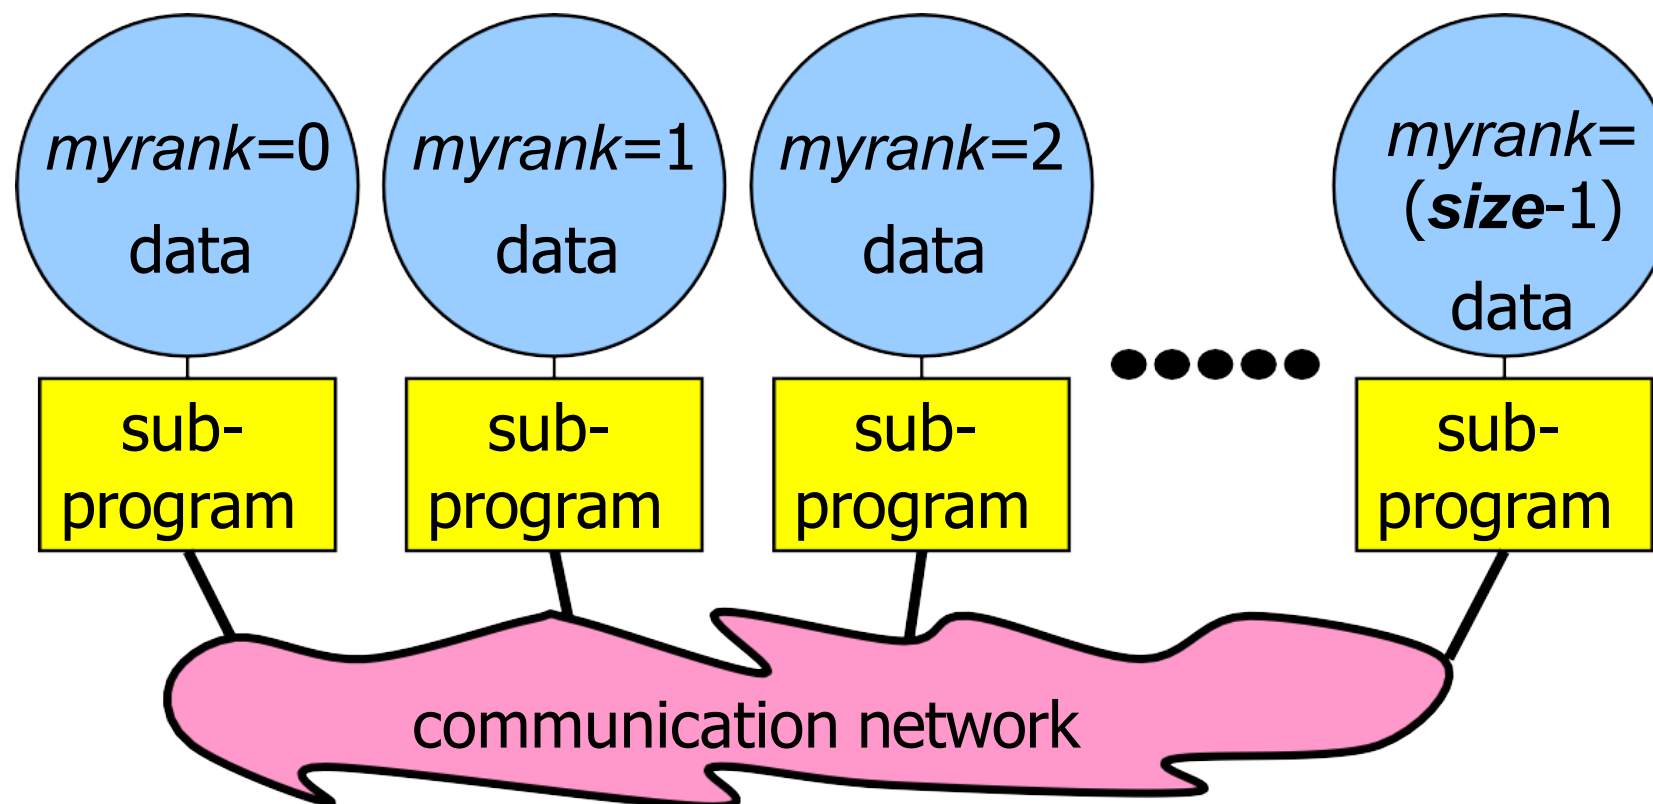
# The Message-Passing Programming Paradigm

- Each processor in a message passing program runs a ***sub-program***:
  - written in a conventional sequential language, e.g., C, Fortran, or Python
  - typically the same on each processor (SPMD),
  - the variables of each sub-program have
    - the same name
    - but different locations (distributed memory) and different data!
    - i.e., all variables are private
  - communicate via special send & receive routines (***message passing***)



# Data and Work Distribution

- the value of *myrank* is returned by special library routine
- the system of *size* processes is started by special MPI initialization program (mpirun or mpiexec)
- all distribution decisions are based on *myrank*
- i.e., which process works on which data





```
#include <stdio.h>
#include <mpi.h>
int main(int argc, char *argv[])
```

Compiled, e.g., with: `mpicc first-example.c`  
 Started, e.g., with: `mpiexec -n 4 ./a.out`  
**Then, this code is running 4 times in parallel !**

```
{ int n;    double result;
  int my_rank, num_procs;
```

application-related data

MPI-related data

```
  MPI_Init(&argc, &argv);
  MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
  MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
```

Now, each process knows who it is:  
 number *my\_rank* out of *num\_procs* processes

```
  if (my_rank == 0)
  { printf("Enter the number of elements (n): \n");
    scanf("%d",&n);
```

reading the application data *n* from stdin only by  
 process 0

```
  }
  MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
```

process 0 is sender, all other  
 processes are receivers

broadcasting the content of variable *n* in process 0  
 into variables *n* in all other processes

```
  result = 1.0 * my_rank * n;
  printf("I am process %i out of %i handling the %ith part of n=%i elements,result=%f\n",
        my_rank, num_procs, my_rank, n, result);
```

doing some **application work** in each process

```
  if (my_rank != 0)
  { MPI_Send(&result,1,MPI_DOUBLE,0,99,MPI_COMM_WORLD);
```

send to process 0

sending some results from  
 all processes (except 0) to process 0

```
  else
```

Process 0: receiving all these messages and, e.g., printing them

```
  { int rank;
    printf("I'm proc 0: My own result is %f \n",result);
    for (rank=1; rank<num_procs; rank++)
    {
      MPI_Recv(&result,1,MPI_DOUBLE,rank,99,
        MPI_COMM_WORLD, MPI_STATUS_IGNORE);
      printf("I'm proc 0: received result of
        process %i is %f \n", rank, result);
    }
  }
```

receiving the message from process *rank*

```
  MPI_Finalize();
```

} Jun Li, Department of Computer Science, CUNY Queens College

Enter the number of elements (n): 100

I am process 0 out of 4 handling the 0th part of n=100 elements, result=0.0  
 I am process 2 out of 4 handling the 2th part of n=100 elements, result=200.0  
 I am process 3 out of 4 handling the 3th part of n=100 elements, result=300.0  
 I am process 1 out of 4 handling the 1th part of n=100 elements, result=100.0  
 I'm proc 0: My own result is 0.0  
 I'm proc 0: received result of process 1 is 100.0  
 I'm proc 0: received result of process 2 is 200.0  
 I'm proc 0: received result of process 3 is 300.0