

```
#include <stdio.h>
#include <mpi.h>
int main(int argc, char *argv[])
```

Compiled, e.g., with: `mpicc first-example.c`
 Started, e.g., with: `mpiexec -n 4 ./a.out`
Then, this code is running 4 times in parallel !

```
{ int n; double result;
  int my_rank, num_procs;
```

application-related data

MPI-related data

```
MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
```

Now, each process knows who it is:
 number *my_rank* out of *num_procs* processes

```
if (my_rank == 0)
{ printf("Enter the number of elements (n): \n");
  scanf("%d",&n);
}
```

reading the application data *n* from stdin only by
 process 0

```
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
```

process 0 is sender, all other
 processes are receivers

broadcasting the content of variable *n* in process 0
 into variables *n* in all other processes

```
result = 1.0 * my_rank * n;
printf("I am process %i out of %i handling the %ith part of n=%i elements,result=%f\n",
       my_rank, num_procs, my_rank, n, result);
```

doing some **application work** in each process

```
if (my_rank != 0)
```

send to process 0

```
{ MPI_Send(&result,1,MPI_DOUBLE,0,99,MPI_COMM_WORLD);
}
```

sending some results from
 all processes (except 0) to process 0

```
else
```

Process 0: receiving all these messages and, e.g., printing them

```
{ int rank;
  printf("I'm proc 0: My own result is %f \n",result);
  for (rank=1; rank<num_procs; rank++)
  {
    MPI_Recv(&result,1,MPI_DOUBLE,rank,99,
MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("I'm proc 0: received result of
process %i is %f \n", rank, result);
  }
}
```

receiving the message from process *rank*

```
MPI_Finalize();
```

Enter the number of elements (n): 100

I am process 0 out of 4 handling the 0th part of n=100 elements, result=0.0
 I am process 2 out of 4 handling the 2th part of n=100 elements, result=200.0
 I am process 3 out of 4 handling the 3th part of n=100 elements, result=300.0
 I am process 1 out of 4 handling the 1th part of n=100 elements, result=100.0
 I'm proc 0: My own result is 0.0
 I'm proc 0: received result of process 1 is 100.0
 I'm proc 0: received result of process 2 is 200.0
 I'm proc 0: received result of process 3 is 300.0


Run `mpiexec -n 4 python first-example.py`

```
1 from mpi4py import MPI
2
3 # application-related data
4 n = None
5 result = None
6
7 comm_world = MPI.COMM_WORLD
8 # MPI-related data
9 my_rank = comm_world.Get_rank() # or my_rank = MPI.COMM_WORLD.Get_rank()
10 num_procs = comm_world.Get_size() # or ditto ...
11
12 if (my_rank == 0):
13     # reading the application data "n" from stdin only by process 0:
14     n = int(input("Enter the number of elements (n): "))
15
16 # broadcasting the content of variable "n" in process 0
17 # into variables "n" in all other processes:
18 n = comm_world.bcast(n, root=0)
19
20 # doing some application work in each process, e.g.:
21 result = 1.0 * my_rank * n
22 print(f"I am process {my_rank} out of {num_procs} handling the {my_rank}ith part of n={n}
23     elements, result={result}")
24
25 if (my_rank != 0):
26     # sending some results from all processes (except 0) to process 0:
27     comm_world.send(result, dest=0, tag=99)
28 else:
29     # receiving all these messages and, e.g., printing them
30     rank = None
31     print(f"I'm proc 0: My own result is {result}")
32     for rank in range(1, num_procs):
33         result = comm_world.recv(source=rank, tag=99)
34         print(f"I'm proc 0: received result of process {rank} is {result}")
```

Access

- A sub-program needs to be connected to a message passing system
- A message passing system is similar to:
 - mail box
 - phone line
 - fax machine
 - etc.
- MPI:
 - sub-program must be linked with an MPI library
 - sub-program must use include file of this MPI library
 - the total program (i.e., all sub-programs of the program) must be started with the MPI startup tool

Messages




- Messages are packets of data moving between sub-programs
- Necessary information for the message passing system:

- sending process
- source location
- source data type
- source data size

- receiving process
- destination location
- destination data type
- destination buffer size

} i.e., the ranks

} 

→ **basic or derived datatypes**

Addressing


- Messages need to have addresses to be sent to.
- Addresses are similar to:
 - mail addresses
 - phone number
 - fax number
 - etc.
- MPI: addresses are ranks of the MPI processes (sub-programs)

Point-to-Point Communication

- Simplest form of message passing.
- One process sends a message to another.
- Different types of point-to-point communication:
 - synchronous send
 - buffered = asynchronous send


Synchronous Sends

- The sender gets an information that the message is received.
- Analogue to the *beep* or *okay-sheet* of a fax.



Buffered = Asynchronous Sends

- Only know when the message has left.



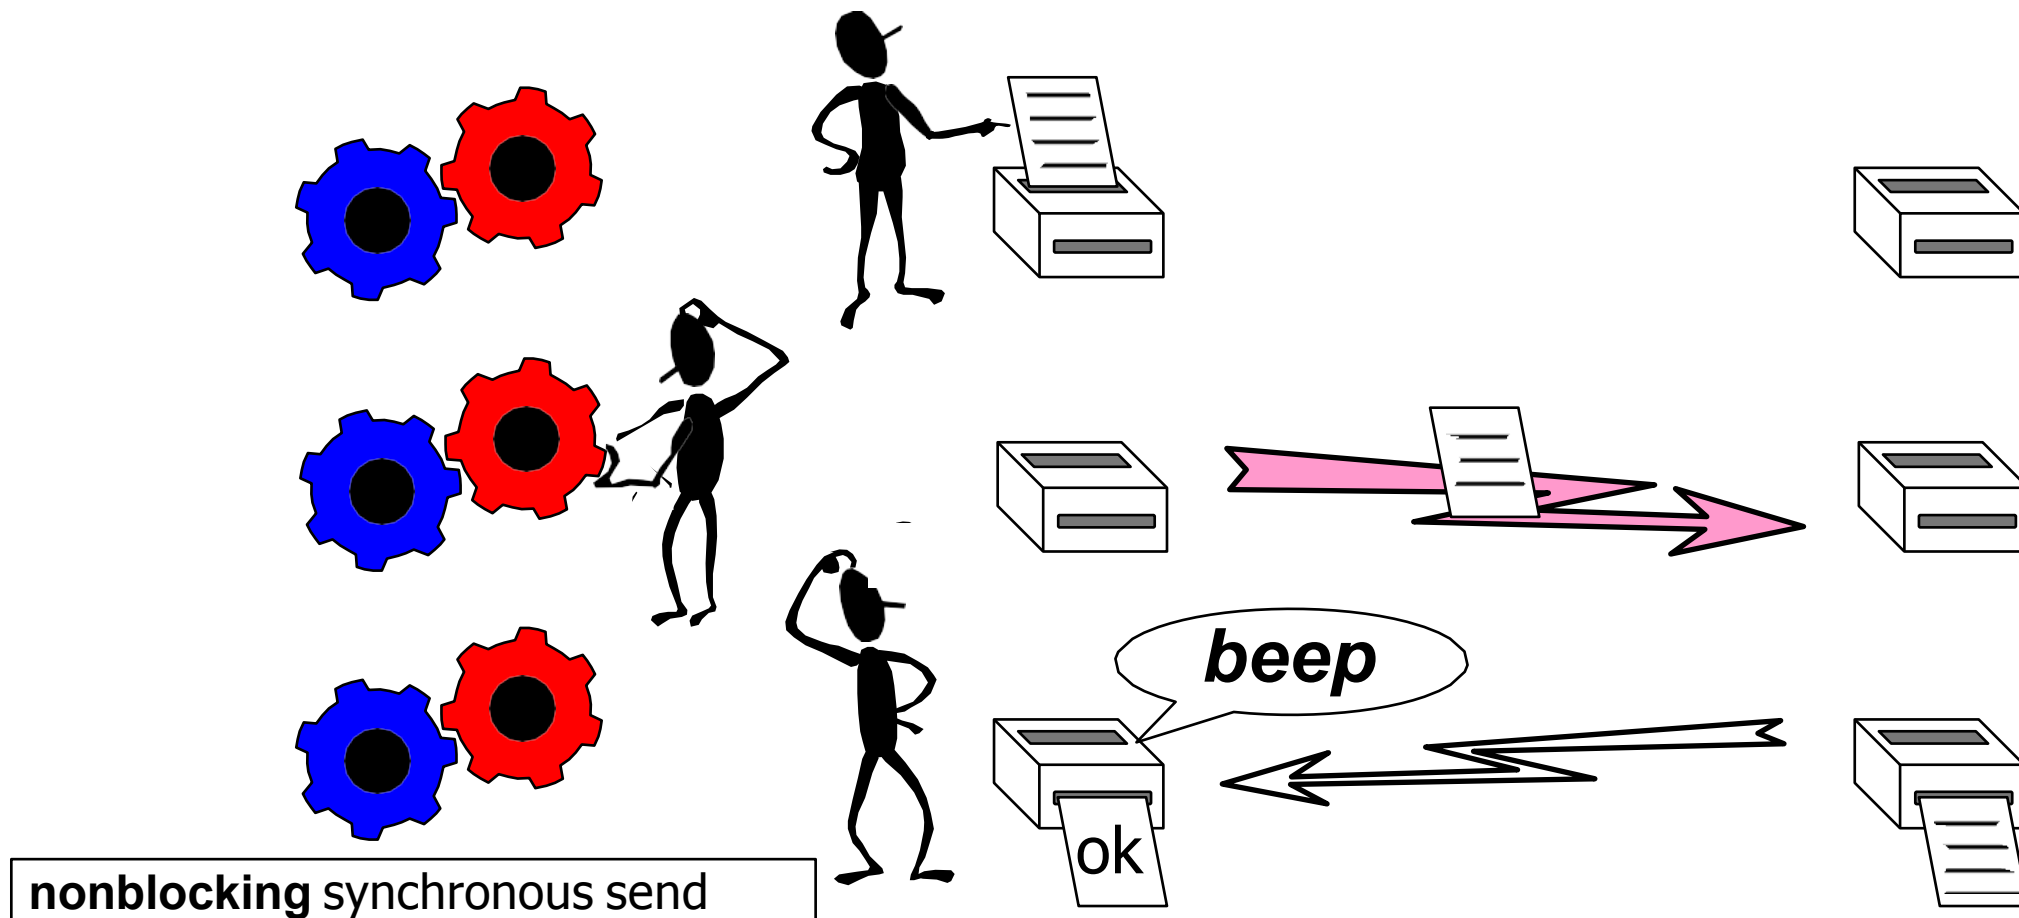
Blocking Operations

- Operations are activities, such as
 - sending (a message)
 - receiving (a message)
- Some operations may **block** until another process acts:
 - synchronous send operation **blocks until** receive is posted;
 - receive operation **blocks until** message was sent.
- Relates to the completion of an operation.
- Blocking subroutine returns only when the operation has completed.

Nonblocking Operations

Nonblocking operations consist of:

- A nonblocking procedure call: it returns immediately and allows the sub-program to perform other work
- At some later time the sub-program must **test** or **wait** for the completion of the nonblocking operation



Non-Blocking Operations (cont'd)


- All nonblocking procedures must have a matching wait (or test) procedure. (Some system or application resources can be freed only when the nonblocking operation is completed.)
- A nonblocking procedure immediately followed by a matching wait is equivalent to a blocking procedure.
- Nonblocking procedures are not the same as sequential subroutine calls:
 - the operation may continue while the application executes the next statements!

Collective Communications

- Collective communication routines are higher level routines.
- Several processes are involved at a time.
- May allow optimized internal implementations, e.g., tree based algorithms.
- Can be built out of point-to-point communications.


Broadcast

- A one-to-many communication.




Reduction Operations

- Combine data from several processes to produce a single result.




Barriers

- Synchronize processes.



Parallel File I/O



calculation on	time for computation	time for <u>serial</u> I/O
1 core	64 min = 98.5 % of total time	1 min = 1.5 % of total time
64 cores	1 min = 50 % of total time	1 min = 50 % of total time

Table: example with serial I/O

-
- 1) waste of resources
 - 2) negative side effects on other users

Process Model & Language Binding

Header files

C

- C / C++

```
#include <mpi.h>
```

Python

- Python

```
from mpi4py import MPI
```

MPI Function Format

In C and Python: case sensitive

C

- C / C++: `error = MPI_Xxxxxx(parameter, ...);`
`MPI_Xxxxxx(parameter, ...);`

Python

- Python: `result_value_or_object = input_mpi_object.mpi_action(parameter, ...)`

direct communication of numPy arrays (like in C)

`comm_world = MPI.COMM_WORLD`
`comm_world.Send((snd_buf, ...), ...)`
`comm_world.Recv((rcv_buf, ...), ...)`

Or with object-serialization:

`comm_world.send(snd_buf, ...)`
`rcv_buf = comm_world.recv(...)`

Mixed cases:

- | | |
|------------------------|----------------------------|
| • MPI procedures in C | <code>MPI_Xxx_mixed</code> |
| • MPI type declaration | <code>MPI_Xxx_mixed</code> |
| • MPI constant | <code>MPI_XXX_UPPER</code> |

Initializing MPI

MPI_Init() must be called before any other MPI routine
(only a few exceptions, e.g., MPI_Initialized)

C

- `int MPI_Init(int *argc, char ***argv)`

MPI-2.0 and higher:
Also
`MPI_Init(NULL, NULL);`

```
#include <mpi.h>
int main(int argc, char **argv)
{
    MPI_Init(&argc, &argv);
    ....
}
```

Python

- `# MPI.Init()`

This call is not needed, because
automatically called at the import of
MPI at the begin of the program

```
from mpi4py import MPI
# MPI.Init() is not needed
....
```

Exiting MPI

C

- C/C++: `int MPI_Finalize()`

Python


- Python: `# MPI.Finalize()`

This call is not needed,
because automatically called at the end of the program

- **Must** be called last by all processes.
- User must ensure the completion of all pending communications (locally) before calling finalize
- After `MPI_Finalize`:
 - Further MPI-calls are forbidden
 - Especially re-initialization with `MPI_Init` is forbidden
 - **May** abort the calling process if its rank in `MPI_COMM_WORLD` is $\neq 0$

Starting the MPI Program


- Start mechanism is implementation dependent
- `mpirun -np number_of_processes ./executable` (most implementations)
- `mpiexec -n number_of_processes ./executable` (with MPI-2 and later)



- The parallel MPI processes exist at least after `MPI_Init` was called.

Communicator MPI_COMM_WORLD

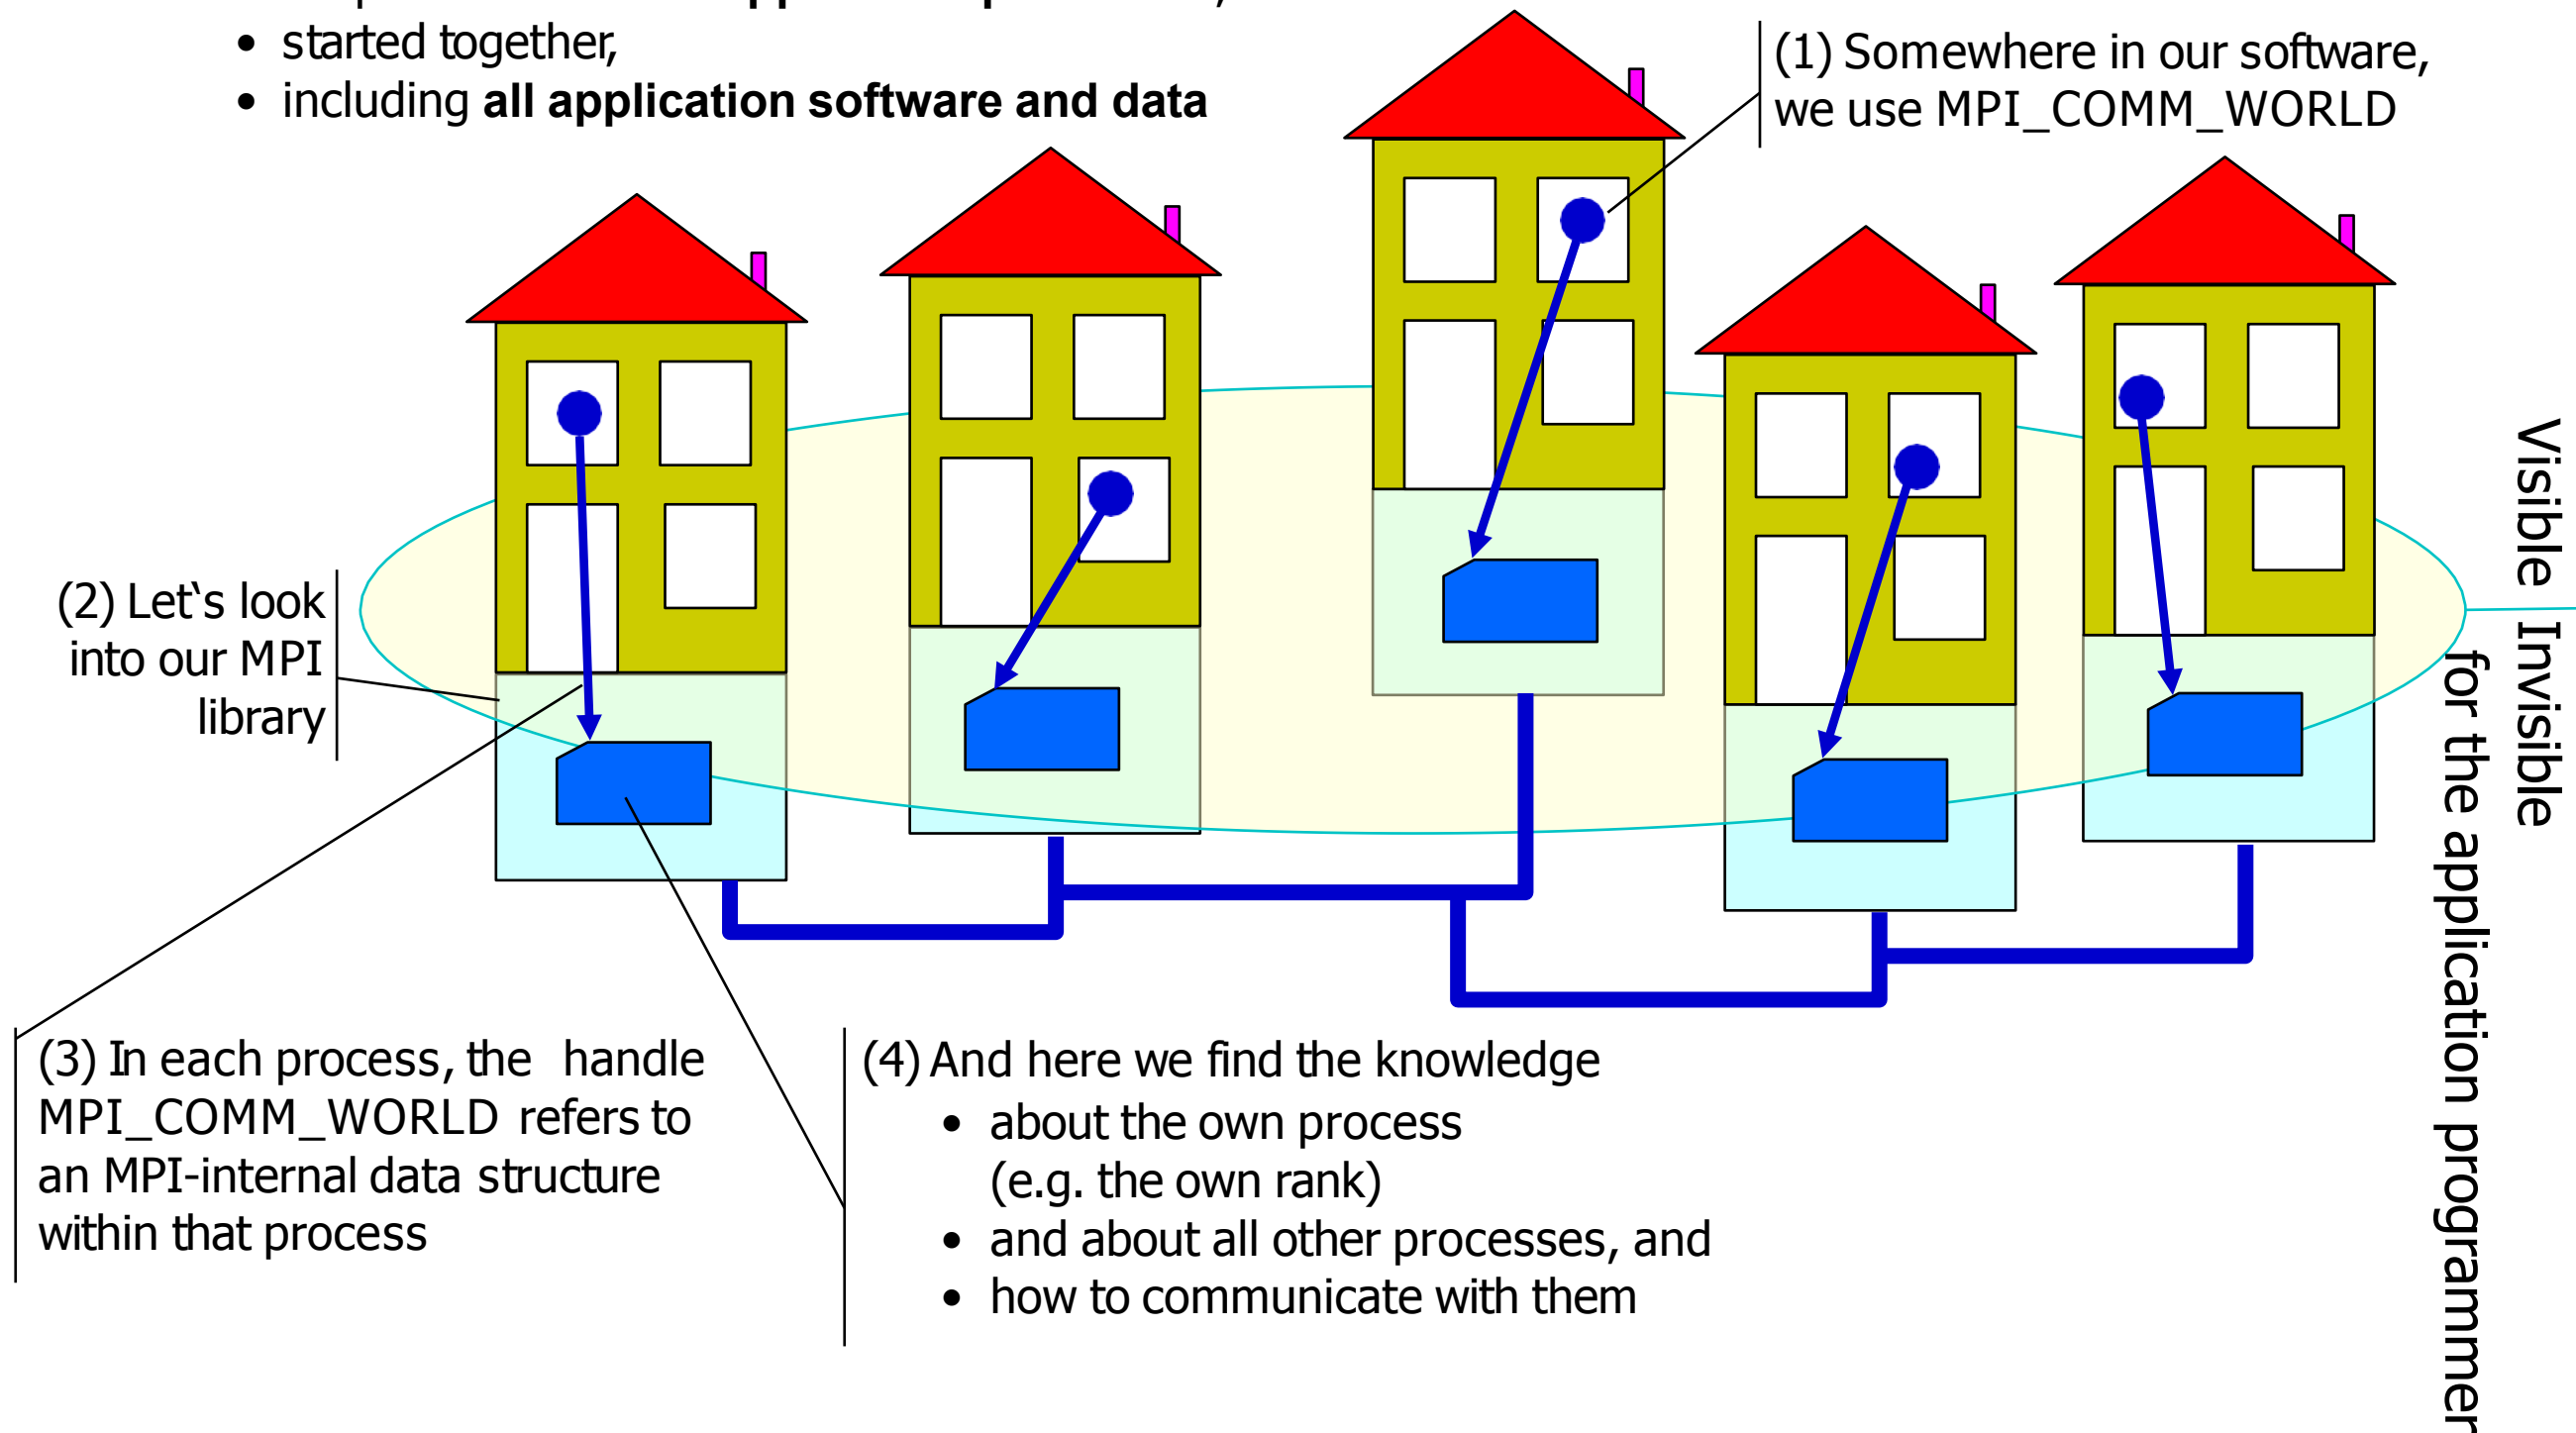
- All processes (= sub-programs) of one MPI program are combined in the **communicator MPI_COMM_WORLD**.
- MPI_COMM_WORLD is a predefined **handle** in
 - mpi.h
 - Each process has its own **rank** in a communicator:
 - starting with 0
 - ending with (size-1)



Handles refer to internal MPI data structures

Example with **five MPI application processes**,

- started together,
- including **all application software and data**



Handles

- Handles identify MPI objects.
- For the programmer, handles are
 - **predefined constants** in C include file mpi.h or **MPI module of mpi4py**
 - Example: MPI_COMM_WORLD or MPI.COMM_WORLD
 - Can be used in initialization expressions or assignments.
 - They are link-time constants, i.e., need not to be compile-time constants.
 - **values returned** by some MPI routines, to be stored in variables, that are defined as
 - C: special MPI typedefs, e.g., MPI_Comm sub_comm;
 - Python: Type of object defined by the creating function, e.g., sub_comm = MPI.COMM_WORLD.Split(...)
- Handles refer to internal MPI data structures

C

Python